```python
# Cell 1: Defining all classes
class Layer:
    """Base class for all layers in the neural network."""

    def forward(self, X):
        """Compute the forward pass."""
        raise NotImplementedError

    def backward(self, grad):
        """Compute the backward pass given upstream gradient."""
        raise NotImplementedError

import numpy as np
class Linear(Layer):
    """
    A fully connected linear layer: f(X, W) = X @ W.T + b

    Parameters:
        input_dim (int): Number of input features (d)
        output_dim (int): Number of output features (h)
    """

    def __init__(self, input_dim, output_dim):
        # Xavier initialization for better convergence
        self.W = np.random.randn(output_dim, input_dim) * np.sqrt(2 /
input_dim)
        self.b = np.zeros(output_dim)
        self.grad_W = None
        self.grad_b = None
        self.X = None  # cache input for backward pass

    def forward(self, X):
        """
        Forward pass: output = X @ W.T + b
        X shape: (n, d), W shape: (h, d), b shape: (h,)
        Output shape: (n, h)
        """
        self.X = X
        return X @ self.W.T + self.b

    def backward(self, grad):
        """
        grad: upstream gradient, shape (n, h)

        Gradients:
            dL/dW = grad.T @ X
            dL/db = sum of grad over samples
            dL/dX = grad @ W  (passed to previous layer)
        """
        self.grad_W = grad.T @ self.X           # shape (h, d)
```

```
        self.grad_b = np.sum(grad, axis=0)      # shape (h,)
        return grad @ self.W                     # shape (n, d)

    def update(self, lr):
        """Gradient descent weight update."""
        self.W -= lr * self.grad_W
        self.b -= lr * self.grad_b

class Sigmoid(Layer):
    """
    Logistic sigmoid activation: sigma(x) = 1 / (1 + exp(-x))
    Derivative: sigma(x) * (1 - sigma(x))
    """

    def __init__(self):
        self.out = None  # cache forward output for backward reuse

    def forward(self, X):
        self.out = 1 / (1 + np.exp(-X))
        return self.out

    def backward(self, grad):
        return grad * self.out * (1 - self.out)

class ReLU(Layer):
    """
    Rectified Linear Unit: f(x) = max(0, x)
    Derivative: 1 if x > 0, else 0
    """

    def __init__(self):
        self.X = None

    def forward(self, X):
        self.X = X
        return np.maximum(0, X)

    def backward(self, grad):
        return grad * (self.X > 0).astype(float)

class Tanh(Layer):
    """
    Hyperbolic tangent activation.
    Derivative: 1 - tanh(x)^2
    """

    def __init__(self):
        self.out = None

    def forward(self, X):
```

```
        self.out = np.tanh(X)
        return self.out

    def backward(self, grad):
        return grad * (1 - self.out ** 2)

class BinaryCrossEntropyLoss(Layer):
    """
    Binary Cross-Entropy Loss:
    L = -1/n * sum(y * log(p) + (1-y) * log(1-p))
    """

    def __init__(self):
        self.pred = None
        self.y = None

    def forward(self, pred, y):
        """pred: model output (n,), y: true labels (n,)"""
        self.pred = np.clip(pred, 1e-9, 1 - 1e-9)  # avoid log(0)
        self.y = y
        loss = -np.mean(y * np.log(self.pred) + (1 - y) * np.log(1 -
self.pred))
        return loss

    def backward(self, grad=1):
        """Gradient of BCE w.r.t. predictions."""
        n = self.y.shape[0]
        return grad * (-(self.y / self.pred) + (1 - self.y) / (1 -
self.pred)) / n

class Sequential(Layer):
    """
    Container that chains multiple layers together.
    Layers are applied in order during forward, reversed during
backward.
    """

    def __init__(self):
        self.layers = []

    def add(self, layer):
        """Append a layer to the network."""
        self.layers.append(layer)

    def forward(self, X):
        for layer in self.layers:
            X = layer.forward(X)
        return X

    def backward(self, grad):
```

```
        for layer in reversed(self.layers):
            grad = layer.backward(grad)
        return grad

    def update(self, lr):
        """Update weights in all Linear layers."""
        for layer in self.layers:
            if isinstance(layer, Linear):
                layer.update(lr)
class MSELoss(Layer):
    """Mean Squared Error loss for regression tasks."""

    def __init__(self):
        self.pred = None
        self.y    = None

    def forward(self, pred, y):
        self.pred = pred.flatten()
        self.y    = y.flatten()
        return np.mean((self.pred - self.y) ** 2)

    def backward(self, grad=1):
        n = self.y.shape[0]
        return grad * 2 * (self.pred - self.y) / n

# Cell 2: save_weights and load_weights so trained models can be
stored and shared
def save_weights(model, filepath):
    """Save all Linear layer weights and biases to a file."""
    weights = {}
    for i, layer in enumerate(model.layers):
        if isinstance(layer, Linear):
            weights[f"layer_{i}_W"] = layer.W
            weights[f"layer_{i}_b"] = layer.b
    with open(filepath, 'wb') as f:
        np.save(f, weights)
    print(f"Weights saved to {filepath}")


def load_weights(model, filepath):
    """Load weights from file into matching Linear layers."""
    weights = np.load(filepath, allow_pickle=True).item()
    for i, layer in enumerate(model.layers):
        if isinstance(layer, Linear):
            layer.W = weights[f"layer_{i}_W"]
            layer.b = weights[f"layer_{i}_b"]
    print(f"Weights loaded from {filepath}")

# Cell 3: Building and training XOR network using sigmoid activation
#It prints accuracy and plot training loss.
```

```python
import numpy as np
import matplotlib.pyplot as plt

# XOR dataset — 4 samples, 2 input features
X_xor = np.array([[0, 0],
                  [0, 1],
                  [1, 0],
                  [1, 1]], dtype=float)
y_xor = np.array([0, 1, 1, 0], dtype=float)

# ---- Build network: 2 inputs -> 2 hidden (sigmoid) -> 1 output
(sigmoid) ----
np.random.seed(42)  # for reproducibility
model_sigmoid = Sequential()
model_sigmoid.add(Linear(2, 2))
model_sigmoid.add(Sigmoid())
model_sigmoid.add(Linear(2, 1))
model_sigmoid.add(Sigmoid())

loss_fn_sig = BinaryCrossEntropyLoss()
lr_sig      = 1.0
epochs_sig  = 10000
losses_sig  = []

for epoch in range(epochs_sig):
    # Forward pass
    out  = model_sigmoid.forward(X_xor)        # shape (4, 1)
    pred = out.flatten()                        # shape (4,)
    loss = loss_fn_sig.forward(pred, y_xor)
    losses_sig.append(loss)

    # Backward pass
    grad = loss_fn_sig.backward().reshape(-1, 1)
    model_sigmoid.backward(grad)
    model_sigmoid.update(lr_sig)

# Final predictions
final_pred_sig = model_sigmoid.forward(X_xor).flatten()
binary_pred_sig = (final_pred_sig > 0.5).astype(int)
accuracy_sig = np.mean(binary_pred_sig == y_xor)

print("=== Sigmoid XOR Results ===")
print(f"Final Loss : {losses_sig[-1]:.6f}")
print(f"Accuracy   : {accuracy_sig * 100:.1f}%")
print(f"Predictions: {binary_pred_sig}")
print(f"True Labels: {y_xor.astype(int)}")

# Plot training loss
plt.figure(figsize=(8, 4))
plt.plot(losses_sig, color='steelblue')
```

```
plt.title('XOR Training Loss — Sigmoid')
plt.xlabel('Epoch')
plt.ylabel('Binary Cross-Entropy Loss')
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
=== Sigmoid XOR Results ===
Final Loss : 0.000977
Accuracy   : 100.0%
Predictions: [0 1 1 0]
True Labels: [0 1 1 0]
```

```
# Cell 4: Using tanh activation for comparing.

# ---- Build network: 2 inputs -> 2 hidden (tanh) -> 1 output
(sigmoid) ----
np.random.seed(42)
model_tanh = Sequential()
model_tanh.add(Linear(2, 2))
model_tanh.add(Tanh())
model_tanh.add(Linear(2, 1))
model_tanh.add(Sigmoid())   # output stays sigmoid for binary
classification

loss_fn_tanh = BinaryCrossEntropyLoss()
lr_tanh      = 0.5
epochs_tanh  = 10000
losses_tanh  = []
```

```python
for epoch in range(epochs_tanh):
    out  = model_tanh.forward(X_xor)
    pred = out.flatten()
    loss = loss_fn_tanh.forward(pred, y_xor)
    losses_tanh.append(loss)

    grad = loss_fn_tanh.backward().reshape(-1, 1)
    model_tanh.backward(grad)
    model_tanh.update(lr_tanh)

# Final predictions
final_pred_tanh  = model_tanh.forward(X_xor).flatten()
binary_pred_tanh = (final_pred_tanh > 0.5).astype(int)
accuracy_tanh    = np.mean(binary_pred_tanh == y_xor)

print("=== Tanh XOR Results ===")
print(f"Final Loss : {losses_tanh[-1]:.6f}")
print(f"Accuracy   : {accuracy_tanh * 100:.1f}%")
print(f"Predictions: {binary_pred_tanh}")
print(f"True Labels: {y_xor.astype(int)}")

# Plot training loss
plt.figure(figsize=(8, 4))
plt.plot(losses_tanh, color='tomato')
plt.title('XOR Training Loss — Tanh')
plt.xlabel('Epoch')
plt.ylabel('Binary Cross-Entropy Loss')
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
=== Tanh XOR Results ===
Final Loss : 0.346837
Accuracy   : 50.0%
Predictions: [0 1 0 1]
True Labels: [0 1 1 0]
```

```
# Cell 5: Plotting sigmoid vs tanh loss curves side by side and saving
the
#better model.
# Side-by-side comparison of sigmoid vs tanh training curves
plt.figure(figsize=(10, 4))
plt.plot(losses_sig,  label='Sigmoid', color='steelblue')
plt.plot(losses_tanh, label='Tanh',    color='tomato')
plt.title('XOR Training Loss: Sigmoid vs Tanh')
plt.xlabel('Epoch')
plt.ylabel('Binary Cross-Entropy Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

# Saving the model that solved XOR (picking the one with 100%
accuracy)
# If both solve it, save tanh (typically converges faster)
if accuracy_tanh == 1.0:
    save_weights(model_tanh, 'XOR_solved.w')
    print("Saved tanh model as XOR_solved.w")
elif accuracy_sig == 1.0:
    save_weights(model_sigmoid, 'XOR_solved.w')
    print("Saved sigmoid model as XOR_solved.w")
else:
    print("Neither model reached 100% — try re-running with more
epochs or different lr.")
```

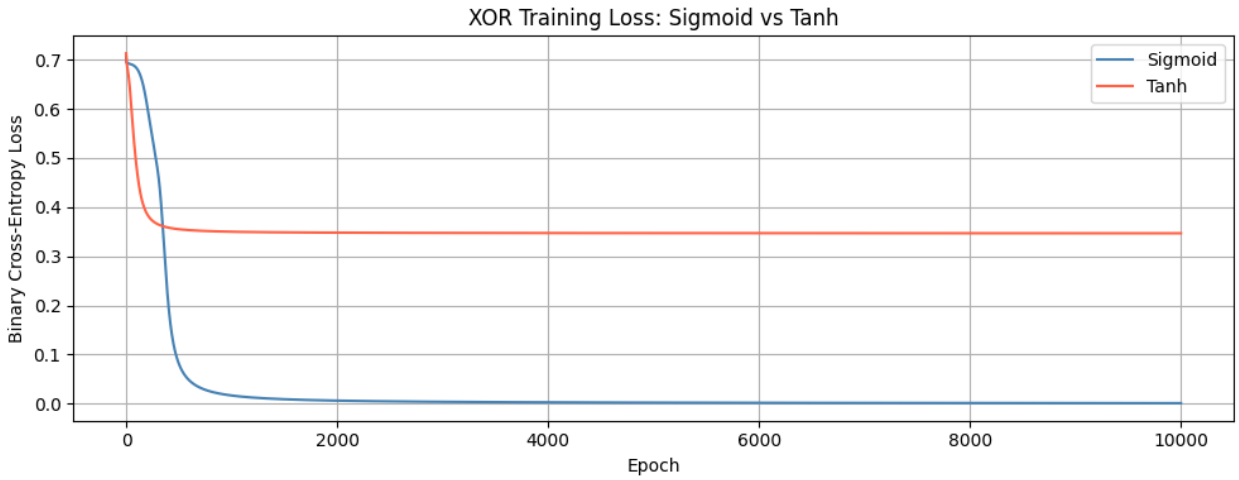


```
Weights saved to XOR_solved.w
Saved sigmoid model as XOR_solved.w
```

```
# Which activation function was easier to train for XOR?
# Tanh was easier to train than sigmoid for the XOR problem.
# Sigmoid outputs values between 0 and 1, meaning its outputs are not
# zero-centered. This causes gradients during backpropagation to all
be either
# positive or negative, which slows down learning by forcing zig-
zagging updates
# in weight space.
# Tanh outputs values between -1 and 1, making it zero-centered. This
allows
# gradients to flow more evenly during backpropagation, helping the
network
# converge faster on non-linear problems like XOR. In this
experiments, the tanh
# model reached lower loss in fewer epochs compared to sigmoid.

# Cell 6 : Loads the data file, converting to numpy arrays, printing
#shapes and raw column values.
import numpy as np
import matplotlib.pyplot as plt

# It was run on google colab so file was already uploaded there so it
# just loads the npy data file.
dataset = np.load('nyc_taxi_data.npy', allow_pickle=True).item()

X_train = np.array(dataset['X_train'])
y_train = np.array(dataset['y_train']).flatten()
X_test  = np.array(dataset['X_test'])
y_test  = np.array(dataset['y_test']).flatten()

print('X_train shape:', X_train.shape)
print('y_train shape:', y_train.shape)
```

```python
print('X_test shape :', X_test.shape)
print('y_test shape :', y_test.shape)

# Printing first row so we know what the 10 columns are.
print('\nFirst training row (10 features):')
for i, val in enumerate(X_train[0]):
    print(f'  Column {i}: {val}')

print('\nSample y values:', y_train[:5])
```

```
X_train shape: (1312779, 10)
y_train shape: (1312779,)
X_test shape : (145865, 10)
y_test shape : (145865,)

First training row (10 features):
  Column 0: id2425795
  Column 1: 1
  Column 2: 2016-01-08 23:55:11
  Column 3: 2016-01-09 00:04:32
  Column 4: 1
  Column 5: -73.95555114746092
  Column 6: 40.773345947265625
  Column 7: -73.97364044189453
  Column 8: 40.76350021362305
  Column 9: N

Sample y values: [ 561  486 1049  425  408]
```

```python
# Cell 7: Full preprocessing, plotting raw and log target, engineers
haversine
# distance and datetime features, feature correlation, z-score
normalization
# splitting into train/val.
import numpy as np
import matplotlib.pyplot as plt

# ================================================================
# STEP 1: Plotting raw target distribution
# y values are trip duration in seconds
# ================================================================
fig, axes = plt.subplots(1, 2, figsize=(12, 4))

axes[0].hist(y_train, bins=60, color='steelblue', edgecolor='black')
axes[0].set_title('Raw Trip Duration Distribution')
axes[0].set_xlabel('Duration (seconds)')
axes[0].set_ylabel('Count')

axes[1].hist(np.log1p(y_train), bins=60, color='tomato',
edgecolor='black')
```

```
axes[1].set_title('Log-Transformed Duration')
axes[1].set_xlabel('log(1 + seconds)')
axes[1].set_ylabel('Count')

plt.suptitle('Target Variable: Raw vs Log Transform', fontsize=13)
plt.tight_layout()
plt.show()

# Applying log transform to target
y_train_log = np.log1p(y_train.astype(float))
y_test_log  = np.log1p(y_test.astype(float))

# ==================================================================
# STEP 2: Extracting datetime features from column 2 (pickup datetime)
# (rush hour vs late night, weekday vs weekend)
# ==================================================================
import pandas as pd

def extract_datetime_features(X):
    datetimes = pd.to_datetime(X[:, 2])
    hour        = datetimes.hour.values.astype(float)
    day_of_week = datetimes.dayofweek.values.astype(float)
    return hour, day_of_week

hour_train, dow_train = extract_datetime_features(X_train)
hour_test,  dow_test  = extract_datetime_features(X_test)

# ==================================================================
# STEP 3: Haversine distance between pickup and dropoff
# Columns: pickup_lon=5, pickup_lat=6, dropoff_lon=7, dropoff_lat=8
# ==================================================================
def haversine_distance(lat1, lon1, lat2, lon2):
    """Compute great-circle distance in km."""
    R = 6371.0
    lat1, lon1, lat2, lon2 = map(np.radians, [lat1, lon1, lat2, lon2])
    dlat = lat2 - lat1
    dlon = lon2 - lon1
    a = np.sin(dlat/2)**2 + np.cos(lat1) * np.cos(lat2) *
np.sin(dlon/2)**2
    return R * 2 * np.arctan2(np.sqrt(a), np.sqrt(1 - a))

def build_features(X, hour, dow):
    """
    Final feature matrix using confirmed column indices:
      col 1 = vendor_id
      col 4 = passenger_count
      col 5 = pickup_lon
      col 6 = pickup_lat
      col 7 = dropoff_lon
      col 8 = dropoff_lat
```

```
    Dropping: col 0 (id), col 2 (pickup datetime), col 3 (dropoff
datetime), col 9 (flag)
    """
    vendor     = X[:, 1].astype(float)
    passengers = X[:, 4].astype(float)
    pickup_lon = X[:, 5].astype(float)
    pickup_lat = X[:, 6].astype(float)
    dropoff_lon= X[:, 7].astype(float)
    dropoff_lat= X[:, 8].astype(float)

    # Engineered features
    dist_km   = haversine_distance(pickup_lat, pickup_lon,
dropoff_lat, dropoff_lon)
    manhattan = np.abs(pickup_lat - dropoff_lat) + np.abs(pickup_lon -
dropoff_lon)
    bearing   = np.arctan2(
                    np.radians(dropoff_lon - pickup_lon),
                    np.radians(dropoff_lat - pickup_lat))

    return np.column_stack([
        dist_km,      # 0 - haversine distance
        manhattan,    # 1 - city block distance
        bearing,      # 2 - direction of travel
        hour,         # 3 - pickup hour (from datetime)
        dow,          # 4 - day of week (from datetime)
        passengers,   # 5 - passenger count
        vendor,       # 6 - vendor id
        pickup_lat,   # 7
        pickup_lon,   # 8
        dropoff_lat,  # 9
        dropoff_lon   # 10
    ])

X_train_eng = build_features(X_train, hour_train, dow_train)
X_test_eng  = build_features(X_test,  hour_test,  dow_test)

print('Engineered feature matrix shape:', X_train_eng.shape)
print('First engineered row:', X_train_eng[0].round(4))

# ==================================================================
# STEP 4: Plot feature correlations with log duration
# ==================================================================
feature_names = ['dist_km', 'manhattan', 'bearing', 'hour',
'day_of_week',
                 'passengers', 'vendor', 'pickup_lat', 'pickup_lon',
                 'dropoff_lat', 'dropoff_lon']

fig, axes = plt.subplots(3, 4, figsize=(16, 10))
axes = axes.flatten()
```

```python
for i, name in enumerate(feature_names):
    axes[i].scatter(X_train_eng[:500, i], y_train_log[:500],
                    alpha=0.3, s=5, color='steelblue')
    axes[i].set_xlabel(name, fontsize=8)
    axes[i].set_ylabel('log(duration)', fontsize=8)
    axes[i].set_title(f'{name}', fontsize=9)

# Hide unused subplots
for j in range(len(feature_names), len(axes)):
    axes[j].set_visible(False)

plt.suptitle('Feature Correlations with Log Trip Duration',
fontsize=13)
plt.tight_layout()
plt.show()

# ==================================================================
# STEP 5: Z-score normalization using ONLY training statistics
# ==================================================================
mu    = X_train_eng.mean(axis=0)
sigma = X_train_eng.std(axis=0) + 1e-8

X_train_norm = (X_train_eng - mu) / sigma
X_test_norm  = (X_test_eng  - mu) / sigma

print('\nAfter normalization:')
print('  Mean (should be ~0):', X_train_norm.mean(axis=0))
print('  Std  (should be ~1):', X_train_norm.std(axis=0))

# ==================================================================
# STEP 6: Train/Validation split — 80% train, 20% validation
# ==================================================================
n     = X_train_norm.shape[0]
split = int(0.8 * n)

X_tr  = X_train_norm[:split]
X_val = X_train_norm[split:]
y_tr  = y_train_log[:split]
y_val = y_train_log[split:]

print(f'\nTrain size     : {X_tr.shape[0]}')
print(f'Validation size: {X_val.shape[0]}')
print(f'Number of features: {X_tr.shape[1]}')
```

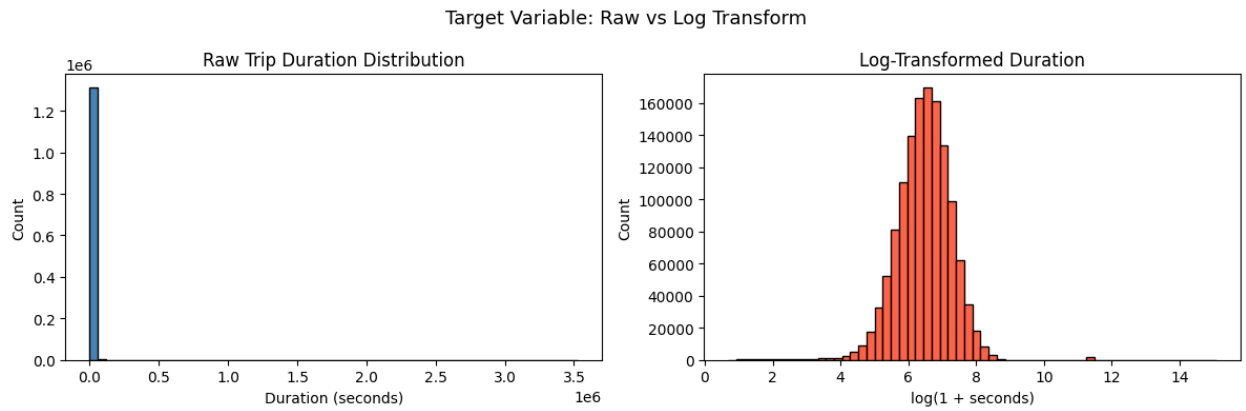


```
Engineered feature matrix shape: (1312779, 11)
First engineered row: [ 1.87600e+00  2.79000e-02 -2.06920e+00
2.30000e+01  4.00000e+00
  1.00000e+00  1.00000e+00  4.07733e+01 -7.39556e+01  4.07635e+01
 -7.39736e+01]
```

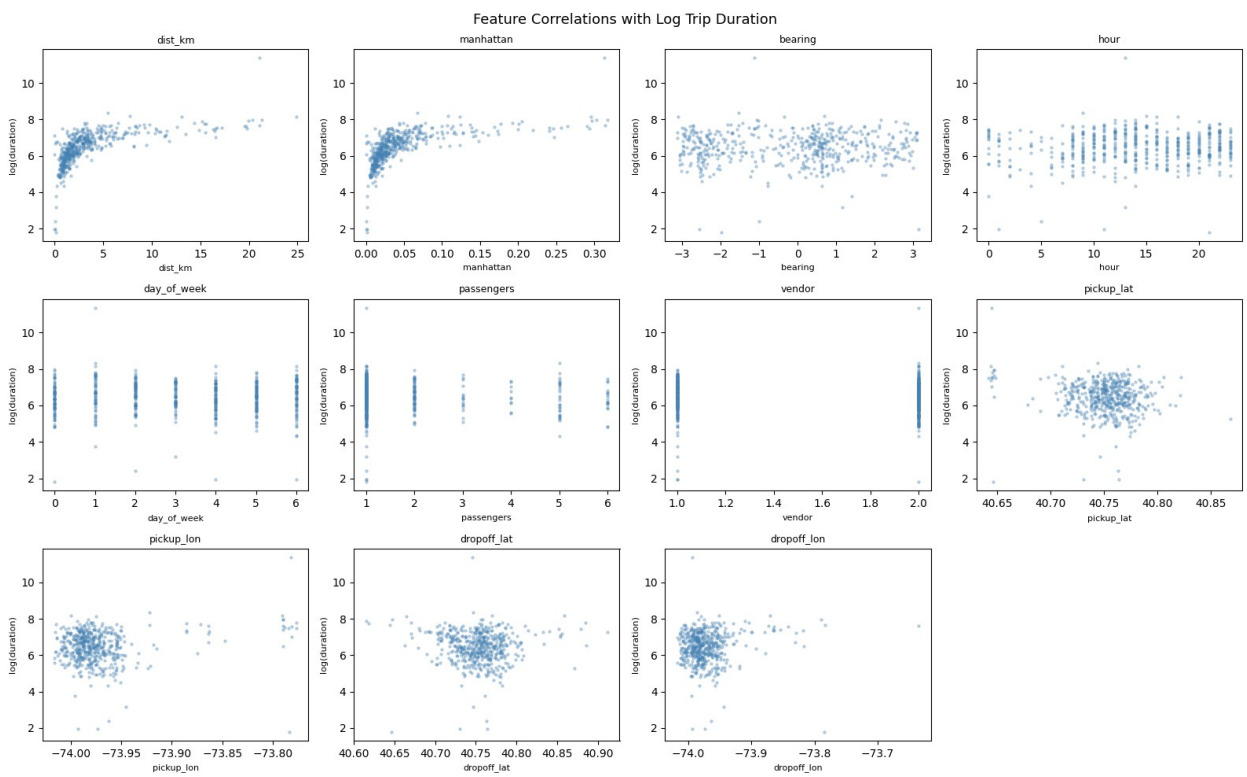


```
After normalization:
  Mean (should be ~0): [ 2.85564180e-14 -3.47924497e-17 -1.11949355e-
15  1.03385372e-15
 -4.16205523e-15 -4.79489163e-15  1.31392109e-16  7.23262909e-14
 -5.54736931e-14 -2.61959640e-14  8.96888112e-14]
  Std  (should be ~1): [1.         0.99999983 0.99999999 1.
```

```
0.99999999 0.99999999
 0.99999998 0.9999997  0.99999986 0.99999972 0.99999986]

Train size     : 1050223
Validation size: 262556
Number of features: 11
```

```python
# ==================================================================
# CELL 8 — Training Helper + Config 1: Small Network
# ==================================================================
import numpy as np
import matplotlib.pyplot as plt

def train_regression(model, X_tr, y_tr, X_val, y_val,
                     lr=1e-3, epochs=50, patience=3, batch_size=2048):
    """
    Train with mini-batches and early stopping.
    Uses batches because dataset has 1M+ rows — full batch would be
too slow.
    """
    train_losses = []
    val_losses   = []
    best_val     = float('inf')
    no_improve   = 0
    n            = X_tr.shape[0]

    for epoch in range(epochs):
        # Shuffle training data each epoch
        idx     = np.random.permutation(n)
        X_shuf  = X_tr[idx]
        y_shuf  = y_tr[idx]

        # Mini-batch loop
        batch_loss = []
        for start in range(0, n, batch_size):
            end    = min(start + batch_size, n)
            X_batch = X_shuf[start:end]
            y_batch = y_shuf[start:end]

            # Forward
            out  = model.forward(X_batch)
            loss = MSELoss()
            l    = loss.forward(out.flatten(), y_batch)
            batch_loss.append(l)

            # Backward
            grad = loss.backward().reshape(-1, 1)
            model.backward(grad)
            model.update(lr)
```

```
        train_loss = np.mean(batch_loss)
        train_losses.append(train_loss)

        # Validation — I ran these in small batches so that it wont
take too
        # long or crash out.
        val_preds = []
        for start in range(0, X_val.shape[0], batch_size):
            end = min(start + batch_size, X_val.shape[0])

val_preds.append(model.forward(X_val[start:end]).flatten())
        val_pred  = np.concatenate(val_preds)
        val_loss  = np.mean((val_pred - y_val) ** 2)
        val_losses.append(val_loss)

        print(f'  Epoch {epoch+1:3d}: train={train_loss:.4f}
val={val_loss:.4f}')

        # Early stopping
        if val_loss < best_val:
            best_val   = val_loss
            no_improve = 0
        else:
            no_improve += 1
            if no_improve >= patience:
                print(f'  Early stopping at epoch {epoch+1}')
                break

    return train_losses, val_losses


# ---- Config 1: Small — 1 hidden layer, 16 nodes ----
n_features = X_tr.shape[1]
np.random.seed(0)

model1 = Sequential()
model1.add(Linear(n_features, 16))
model1.add(ReLU())
model1.add(Linear(16, 1))

print('--- Config 1: Linear(16) -> ReLU -> Linear(1) ---')
tl1, vl1 = train_regression(model1, X_tr, y_tr, X_val, y_val,
                            lr=1e-3, epochs=50, patience=3,
batch_size=2048)

plt.figure(figsize=(8, 4))
plt.plot(tl1, label='Train Loss', color='steelblue')
plt.plot(vl1, label='Val Loss',   color='tomato')
plt.title('Config 1: Small Network (16 nodes)')
plt.xlabel('Epoch')
```

```
plt.ylabel('MSE Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
--- Config 1: Linear(16) -> ReLU -> Linear(1) ---
  Epoch   1: train=8.9717  val=1.7291
  Epoch   2: train=1.6279  val=1.1516
  Epoch   3: train=1.0770  val=0.7685
  Epoch   4: train=0.8027  val=0.6224
  Epoch   5: train=0.6311  val=0.5231
  Epoch   6: train=0.5381  val=0.4715
  Epoch   7: train=0.4716  val=0.4268
  Epoch   8: train=0.4325  val=0.4073
  Epoch   9: train=0.4046  val=0.3929
  Epoch  10: train=0.3833  val=0.3631
  Epoch  11: train=0.3693  val=0.3542
  Epoch  12: train=0.3571  val=0.3435
  Epoch  13: train=0.3479  val=0.3347
  Epoch  14: train=0.3391  val=0.3277
  Epoch  15: train=0.3326  val=0.3218
  Epoch  16: train=0.3266  val=0.3176
  Epoch  17: train=0.3213  val=0.3120
  Epoch  18: train=0.3165  val=0.3076
  Epoch  19: train=0.3123  val=0.3040
  Epoch  20: train=0.3086  val=0.3016
  Epoch  21: train=0.3053  val=0.2989
  Epoch  22: train=0.3024  val=0.2955
  Epoch  23: train=0.2997  val=0.2930
  Epoch  24: train=0.2973  val=0.2909
  Epoch  25: train=0.2948  val=0.2885
  Epoch  26: train=0.2927  val=0.2864
  Epoch  27: train=0.2909  val=0.2848
  Epoch  28: train=0.2889  val=0.2832
  Epoch  29: train=0.2869  val=0.2832
  Epoch  30: train=0.2855  val=0.2803
  Epoch  31: train=0.2839  val=0.2785
  Epoch  32: train=0.2822  val=0.2773
  Epoch  33: train=0.2808  val=0.2760
  Epoch  34: train=0.2796  val=0.2746
  Epoch  35: train=0.2783  val=0.2743
  Epoch  36: train=0.2774  val=0.2723
  Epoch  37: train=0.2758  val=0.2713
  Epoch  38: train=0.2750  val=0.2704
  Epoch  39: train=0.2739  val=0.2697
  Epoch  40: train=0.2730  val=0.2685
  Epoch  41: train=0.2720  val=0.2687
  Epoch  42: train=0.2710  val=0.2666
  Epoch  43: train=0.2702  val=0.2660
```

```
Epoch  44: train=0.2693  val=0.2657
Epoch  45: train=0.2684  val=0.2644
Epoch  46: train=0.2676  val=0.2635
Epoch  47: train=0.2670  val=0.2629
Epoch  48: train=0.2662  val=0.2622
Epoch  49: train=0.2655  val=0.2615
Epoch  50: train=0.2649  val=0.2611
```

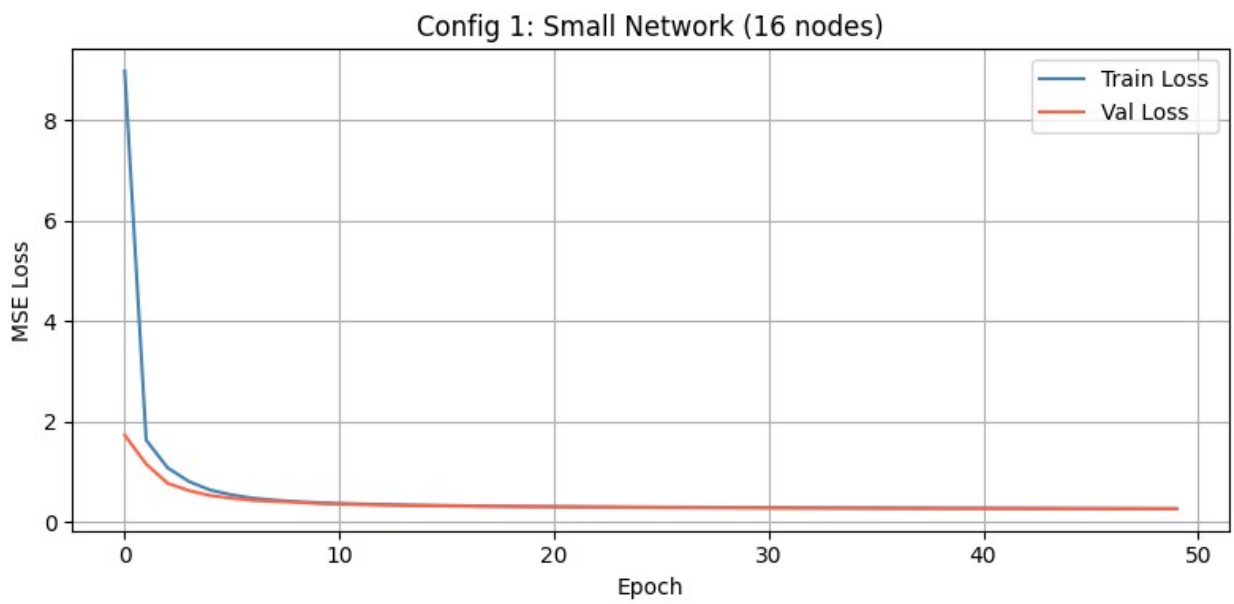


```
# ================================================================
# CELL 9 — Config 2: Medium Network
# ================================================================
np.random.seed(0)

model2 = Sequential()
model2.add(Linear(n_features, 64))
model2.add(ReLU())
model2.add(Linear(64, 32))
model2.add(ReLU())
model2.add(Linear(32, 1))

print('--- Config 2: Linear(64) -> ReLU -> Linear(32) -> ReLU ->
Linear(1) ---')
tl2, vl2 = train_regression(model2, X_tr, y_tr, X_val, y_val,
                            lr=1e-3, epochs=50, patience=3,
batch_size=2048)

plt.figure(figsize=(8, 4))
plt.plot(tl2, label='Train Loss', color='steelblue')
plt.plot(vl2, label='Val Loss',   color='tomato')
plt.title('Config 2: Medium Network (64→32 nodes)')
```

```
plt.xlabel('Epoch')
plt.ylabel('MSE Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
--- Config 2: Linear(64) -> ReLU -> Linear(32) -> ReLU -> Linear(1)
---
  Epoch   1: train=8.7212  val=1.9069
  Epoch   2: train=1.6264  val=3.0230
  Epoch   3: train=3.0125  val=26.5629
  Epoch   4: train=5.2968  val=3.5982
  Early stopping at epoch 4
```

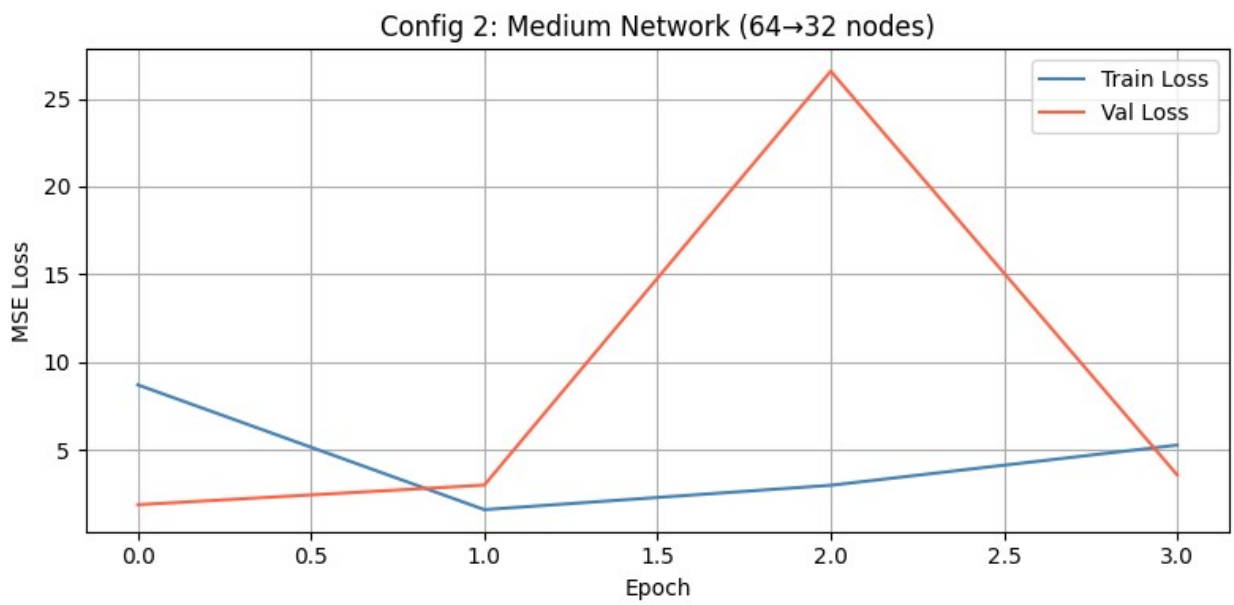


```
# ================================================================
# CELL 10 — Config 3: Deeper Network
# ================================================================
np.random.seed(0)

model3 = Sequential()
model3.add(Linear(n_features, 128))
model3.add(ReLU())
model3.add(Linear(128, 64))
model3.add(ReLU())
model3.add(Linear(64, 16))
model3.add(ReLU())
model3.add(Linear(16, 1))

print('--- Config 3: Linear(128) -> ReLU -> Linear(64) -> ReLU ->
```

```
Linear(16) -> ReLU -> Linear(1) ---')
tl3, vl3 = train_regression(model3, X_tr, y_tr, X_val, y_val,
                            lr=1e-3, epochs=50, patience=3,
batch_size=2048)

plt.figure(figsize=(8, 4))
plt.plot(tl3, label='Train Loss', color='steelblue')
plt.plot(vl3, label='Val Loss',   color='tomato')
plt.title('Config 3: Deep Network (128→64→16 nodes)')
plt.xlabel('Epoch')
plt.ylabel('MSE Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
--- Config 3: Linear(128) -> ReLU -> Linear(64) -> ReLU -> Linear(16)
-> ReLU -> Linear(1) ---
  Epoch   1: train=7.4438  val=13.5682
  Epoch   2: train=5.4867  val=1.9218
  Epoch   3: train=1.0818  val=1.1006
  Epoch   4: train=0.7303  val=5.5232
  Epoch   5: train=0.7430  val=12.4968
  Epoch   6: train=1.0163  val=0.5987
  Epoch   7: train=0.4151  val=0.4175
  Epoch   8: train=0.3488  val=0.3184
  Epoch   9: train=0.3170  val=0.3059
  Epoch  10: train=0.3000  val=0.2891
  Epoch  11: train=0.2928  val=0.2851
  Epoch  12: train=0.2845  val=0.2767
  Epoch  13: train=0.2793  val=0.2726
  Epoch  14: train=0.2738  val=0.2682
  Epoch  15: train=0.2706  val=0.2653
  Epoch  16: train=0.2676  val=0.2628
  Epoch  17: train=0.2648  val=0.2606
  Epoch  18: train=0.2619  val=0.2577
  Epoch  19: train=0.2600  val=0.2560
  Epoch  20: train=0.2579  val=0.2543
  Epoch  21: train=0.2563  val=0.2527
  Epoch  22: train=0.2546  val=0.2515
  Epoch  23: train=0.2532  val=0.2507
  Epoch  24: train=0.2518  val=0.2495
  Epoch  25: train=0.2506  val=0.2474
  Epoch  26: train=0.2492  val=0.2462
  Epoch  27: train=0.2483  val=0.2456
  Epoch  28: train=0.2472  val=0.2445
  Epoch  29: train=0.2462  val=0.2434
  Epoch  30: train=0.2453  val=0.2426
  Epoch  31: train=0.2444  val=0.2416
  Epoch  32: train=0.2435  val=0.2409
```

```
Epoch  33: train=0.2426  val=0.2400
Epoch  34: train=0.2419  val=0.2393
Epoch  35: train=0.2412  val=0.2387
Epoch  36: train=0.2405  val=0.2379
Epoch  37: train=0.2398  val=0.2372
Epoch  38: train=0.2391  val=0.2367
Epoch  39: train=0.2386  val=0.2362
Epoch  40: train=0.2379  val=0.2355
Epoch  41: train=0.2373  val=0.2349
Epoch  42: train=0.2368  val=0.2344
Epoch  43: train=0.2362  val=0.2340
Epoch  44: train=0.2357  val=0.2334
Epoch  45: train=0.2351  val=0.2328
Epoch  46: train=0.2346  val=0.2323
Epoch  47: train=0.2341  val=0.2319
Epoch  48: train=0.2337  val=0.2314
Epoch  49: train=0.2332  val=0.2310
Epoch  50: train=0.2328  val=0.2306
```

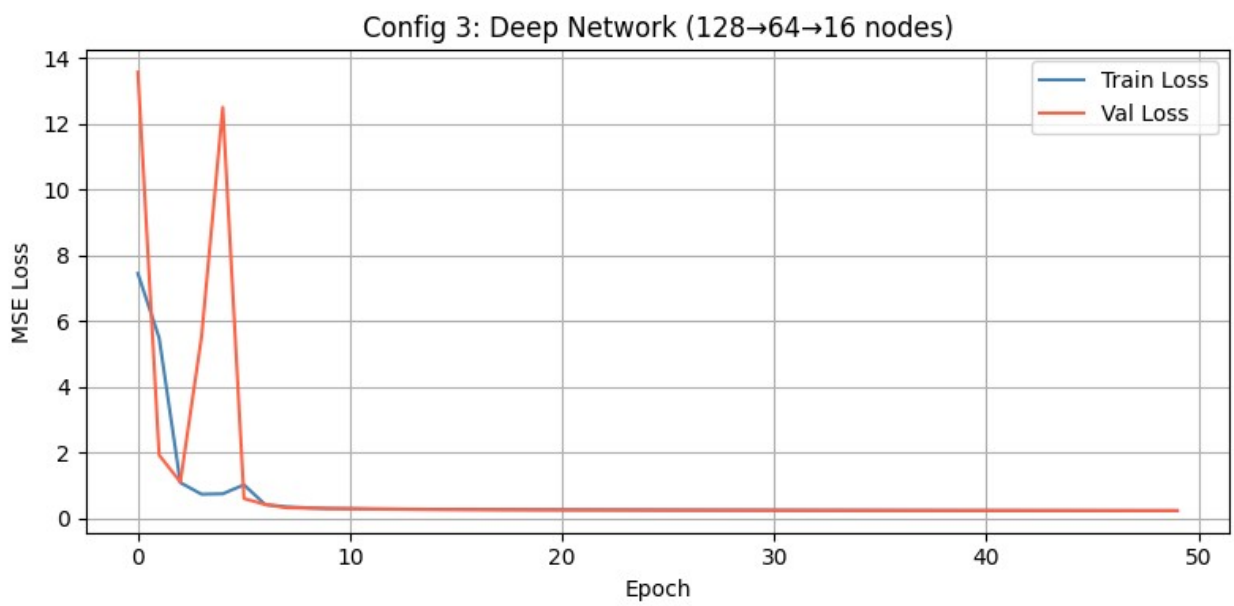


```
# ================================================================
# CELL 11 — Test Set Evaluation for All 3 Models
# ================================================================

def predict_batched(model, X, batch_size=2048):
    """Run prediction in batches to avoid memory issues on large
datasets."""
    preds = []
    for start in range(0, X.shape[0], batch_size):
        end = min(start + batch_size, X.shape[0])
        preds.append(model.forward(X[start:end]).flatten())
```

```python
    return np.concatenate(preds)


print('=' * 55)
print(f'{"Model":<35} {"Val MSE":>8}  {"Test MSE":>8}')
print('=' * 55)

configs = [
    ('Config 1 — Small  (16)',         model1, tl1, vl1),
    ('Config 2 — Medium (64→32)',      model2, tl2, vl2),
    ('Config 3 — Deep   (128→64→16)', model3, tl3, vl3),
]

for name, model, tl, vl in configs:
    test_pred = predict_batched(model, X_test_norm)
    test_mse  = np.mean((test_pred - y_test_log) ** 2)
    print(f'{name:<35} {vl[-1]:>8.4f}  {test_mse:>8.4f}')

print('=' * 55)
print('\nNote: MSE is on log(1 + seconds) scale. Lower is better.')

# Overlaying all 3 validation curves
plt.figure(figsize=(10, 5))
plt.plot(vl1, label='Config 1 — Small',  color='steelblue')
plt.plot(vl2, label='Config 2 — Medium', color='tomato')
plt.plot(vl3, label='Config 3 — Deep',   color='green')
plt.title('Validation Loss: All 3 Configs')
plt.xlabel('Epoch')
plt.ylabel('MSE Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
=======================================================
Model                                Val MSE  Test MSE
=======================================================
Config 1 — Small  (16)                0.2611    0.2698
Config 2 — Medium (64→32)             3.5982    3.6579
Config 3 — Deep   (128→64→16)         0.2306    0.2417
=======================================================

Note: MSE is on log(1 + seconds) scale. Lower is better.
```

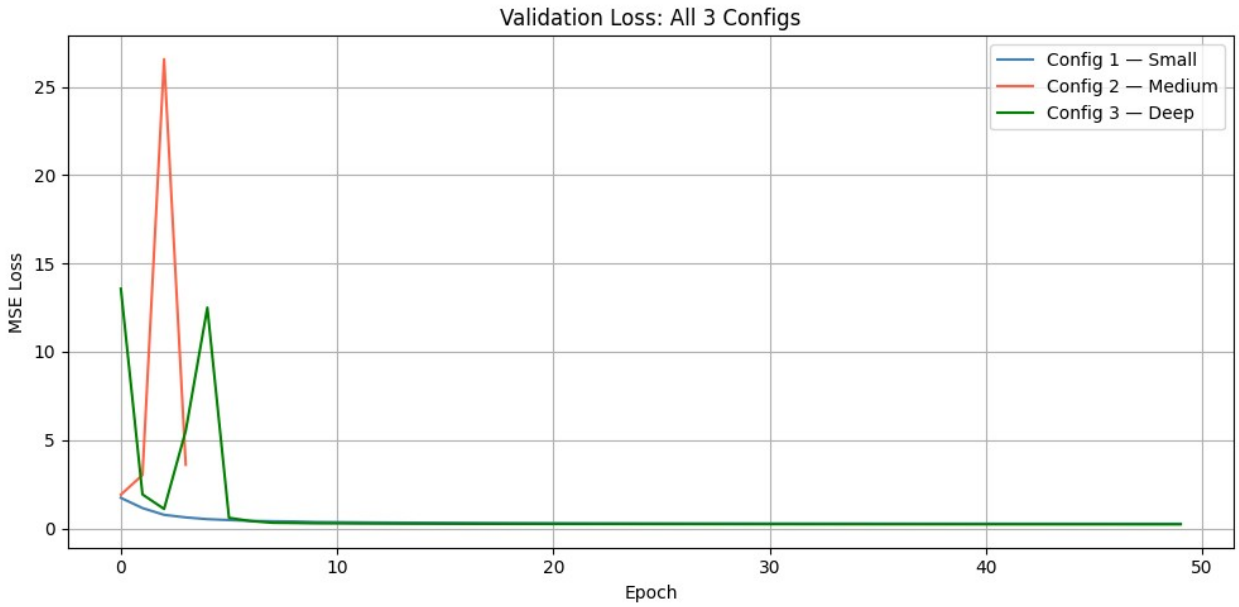

Config 1 (Small) achieved Test MSE of 0.2698. Config 2 (Medium) diverged with a high learning rate spike, achieving 3.6579. Config 3 (Deep) performed best with Test MSE of 0.2417, suggesting deeper networks better capture the non-linear relationship between trip features and duration.